

SETHL: PARTIALLY SPECIFIED HAND ARTICULATION CONTROL FOR VIDEO DIFFUSION

Haotian Lu Xiao-Ping Zhang*

Tsinghua University
haotialu666@gmail.com, xpzhang@ieee.org

ABSTRACT

Pose-conditioned video generators assume complete trajectories, although users may specify only important fingers or intervals. We formulate *partially specified hand control*: obey declared constraints while preserving variation elsewhere. SetHL extends Hand Labanotation (HL) from fixed symbols to distributions over unspecified cells. A masked transformer completes missing symbols, a certified in-cell sampler provides continuous variation without changing HL labels, and a kinematic layer drives a Wan2.1-Fun-Control 1.3B model trained with rank-16 LoRA. On quality-filtered WLASL, three-seed experiments show that SetHL improves specified-symbol accuracy over continuous and learned vector-quantized controls by 1.59 and 1.34 points, and reduces leakage by 3.87 and 2.80 points. The primary compliance–diversity hypervolume differences are not significant, while conditional denoising MSE improves by 2.81% versus continuous control. These results support structured partial constraints as an editable video-diffusion interface that separates controlled from free hand articulation.

Index Terms— controllable video generation, hand motion, Hand Labanotation, diffusion models

1. INTRODUCTION

Hand motion is both semantically dense and difficult to synthesize. Many pose-based controllers consume a complete joint trajectory, while sparse-control methods usually sparsify time or landmarks rather than letting authors declare which articulations are unconstrained [1, 2, 3, 4]. Complete trajectories are suitable for reenactment, but form a poor authoring interface: a user may require an index finger to point inward during one interval while leaving the remaining articulation unspecified. Treating missing coordinates as zeros or hallucinating a single completed skeleton confuses “free to vary” with “fixed to an arbitrary estimate.”

Hand Labanotation (HL) [5] provides a useful alternative. It describes each of 20 hand-local bone directions by one of 26 interpretable spatial symbols. The original work used HL for

recognition, documentation, and approximate one-shot skeleton recovery. We ask a different question: can a partial symbolic program define a *set* of valid videos rather than one trajectory?

We propose SetHL-Wan. Specified HL symbols are preserved exactly, whereas a temporal completion network represents each unspecified cell by a categorical posterior. Samples are forward-kinematically spatialized and lifted to the control-latent interface of Wan [6]. Only low-rank adapters [7], the completion network, and the lifting module are trained. Crucially, no dense joint coordinates bypass the symbolic stream.

Our contributions are: (1) a task definition and metric that jointly measures constraint compliance and diversity localized to unconstrained hand parts; (2) a probabilistic extension of HL with certified symbol-preserving continuous samples, coupled to a differentiable kinematic Wan controller; and (3) matched-budget tests against continuous and learned-token alternatives, including structured missing-control and noise stress tests.

2. RELATED WORK

Controllable video diffusion. ControlNet [1] established spatial conditioning for diffusion models, while SparseCtrl [2] studied temporally sparse video conditions. Recent work conditions video models on pose or trajectories [3, 4, 8, 9, 10]. Sparse 3-D hand control has also addressed occlusion and geometric fidelity [11], but “sparse” there describes landmarks or missing observations, not author-declared degrees of freedom. These methods seek one accurate motion trace; our target is selective constraint with localized residual variation. Learned 3-D motion and mesh tokens avoid rasterized controls [12, 13], but encode complete motion with task-specific vocabularies rather than an interpretable, partially specified program.

Hand and sign generation. WLASL [14] enables large-scale word-level sign recognition. Recent generators improve sign fidelity through pose adapters [15] or compressed multi-modal tokens [16]. They motivate stronger hand evaluation, but do not expose a symbolic partial-control interface.

Symbolic hand motion. HL [5] quantizes hand-local directions and is robust to small pose noise. We retain its fixed

*Corresponding author.




Fig. 1. SetHL treats a partial hand instruction as a set rather than a single completed skeleton. Given HL cells are copied exactly; a masked transformer predicts distributions for free cells, which are sampled and spatialized by forward kinematics before conditioning a LoRA-adapted Wan model. Continuous within-cell samples preserve their selected HL symbol by construction.

semantics, but replace deterministic cell-center recovery with a posterior over unspecified cells and train it as a generative condition. A learned 26-entry VQ codebook is included to test whether gains come merely from discretization.

3. METHOD

Let $X = (x_1, \dots, x_T)$ be a video and $J_t \in \mathbb{R}^{2 \times 21 \times 3}$ its two-hand landmarks. The control mask $M \in \{0, 1\}^{T \times 2 \times 20}$ indicates which bone-time cells the user specifies.

3.1. Partial Hand-Labanotation programs

For each hand we construct the palm frame R_t from the wrist, index metacarpophalangeal joint, middle metacarpophalangeal joint, and little-finger metacarpophalangeal joint as in [5]. A bone direction $d_{t,b}$ is mapped into this frame and assigned to its nearest member of the fixed 26-direction codebook $C = \{c_k\}_{k=1}^{26}$:

$$z_{t,b} = \arg \max_k \langle R_t^\top d_{t,b}, c_k \rangle. \quad (1)$$

The input program contains $(M, M \odot z)$, wrist trajectories, palm frames, and reference bone lengths. It never contains per-joint image coordinates. Specified cells are one-hot, preventing continuous sub-cell information from leaking through soft labels. We interpret each symbol as its spherical Voronoi cell rather than only its center.

For missing cells, a transformer q_ϕ attends jointly over time, hands, and the 20-bone trees:

$$\begin{aligned} q_{t,b} &= q_\phi(z_{\text{obs}}, M), \\ \mathcal{L}_{\text{HL}} &= -\log q_{t,b}(z_{t,b}) + \lambda_h [1 - \langle \bar{c}_{t,b}, c_{z_{t,b}} \rangle]. \end{aligned} \quad (2)$$

where $\bar{c}$ is the normalized posterior mean. The second term respects the geometry of nearby HL regions instead of treating all symbol mistakes as equally severe.

3.2. Kinematic control-latent lifting

A hard straight-through categorical sample first selects a center c_z . For unit $v \perp c_z$, we draw

$$\tilde{d} = \cos \alpha c_z + \sin \alpha v, \quad \alpha < \frac{1}{2} \min_{j \neq z} \arccos(c_z^\top c_j). \quad (3)$$

By the spherical triangle inequality, c_z remains the nearest codeword. This gives continuous variation with a certificate that the HL instruction is unchanged. Forward kinematics reconstructs bone endpoints from wrist position, palm frame, and bone length. Bone tokens are bilinearly splatted along their 2-D segments and decoded by three small 3-D convolutions into a 16-channel latent $H_\psi(P)$ at Wan’s temporal and spatial resolution. We concatenate it with the first-frame latent, matching the native Wan-Control input contract.

Given clean latent v , noise ϵ , and flow-matching interpolation $v_\sigma = (1 - \sigma)v + \sigma\epsilon$, we optimize

$$\mathcal{L} = \|f_{\theta, \Delta}(v_\sigma, \sigma, H_\psi(P)) - (\epsilon - v)\|_2^2 + 0.1 \mathcal{L}_{\text{HL}}, \quad (4)$$

where the 1.3B Wan backbone θ is frozen and Δ denotes rank-16 LoRA parameters. At inference, independent posterior samples alter the control trajectory only where the program permits it, while diffusion noise supplies appearance variation.

4. EXPERIMENTS

Data and protocol. We use 2,252 WLASL videos with reliable MediaPipe Hands [17] tracks, retaining the official




Fig. 2. Generated-video comparison on the source at the median three-seed SetHL hypervolume (a deterministic selection rule). Teal mask cells are specified and gray cells are free; frames are taken at the center of the withheld interval. Columns show paired diffusion seeds under continuous and SetHL control.

1,579/404/269 train/validation/test partition. Each source video contributes exactly one 17-frame clip; no overlapping window crosses a split. Frames are resized to 256^2 . Every learned method uses the same Wan checkpoint, rank-16 LoRA, 3,000 updates, and three seeds. SetHL trains 21.85M parameters in total (LoRA plus its completion and lifting modules), while the matched continuous model trains 21.84M. With cached Wan-VAE latents, each 3,000-step seed takes about 28 minutes on one Ascend 910B2C NPU. Text and CLIP embeddings are fixed to zero for every method so the comparison isolates first-frame and motion-control conditioning. We select $\lambda_h = 1$ from $\{0, 0.25, 1\}$ using validation data only. We likewise select posterior temperature 1 from $\{1, 1.5, 2\}$ on eight fixed validation sources and share it across all probabilistic categorical controls.

Baselines and metrics. We compare continuous Gaussian directions, a training-set 26-entry VQ codebook, hard HL centers, SetHL without the spherical loss, and full SetHL. Continuous and raster controls retain exact observed geometry and therefore serve as information-rich upper bounds, while parameter counts and training updates remain matched. At inference we also disable within-cell residuals in the trained SetHL model, isolating their contribution without changing learned weights. Training samples masks that remove a distal joint set, one complete finger, or a contiguous temporal interval. Standalone completion averages all three; the primary generated-video benchmark uses the predeclared

interval mask. A secondary robustness policy is selected between finger and distal masks on the fixed eight-source validation subset, then evaluated once on the test set. We report specified-symbol accuracy, specified angular error, hidden-cell angular diversity, specified-cell symbol disagreement, hand detection, temporal acceleration error, and denoising MSE. The primary endpoint is the Pareto hypervolume dominated by compliance-localized-diversity operating points (reference $(0, 0)$) obtained by scaling the control latent. Localized diversity multiplies hidden-cell angular diversity by semantic agreement on specified cells and hand-detection rate. Thus semantic leakage and untrackable hands are penalized while legal angular variation inside an HL cell is not. Confidence intervals use 10,000 paired bootstrap resamples of source videos. Only the predeclared hypervolume test is confirmatory; secondary intervals are pointwise and unadjusted for multiplicity. Generated-video metrics use all 27 test clips meeting the fixed criterion that both hands are tracked in at least 90% of source frames. As an evaluator ceiling, re-encoding these real clips and rerunning the same tracker gives 92.81% hand detection, 78.36% HL agreement, and 10.24° angular error against the archived pseudo-labels. As a post-hoc quality diagnostic, we also measure source-conditioned cosine similarity in the penultimate feature space of a Kinetics-400-pretrained R3D-18 [18]; this metric was not used for model selection.

Main results. Table 1 reports the locked comparison.

Table 1. Interval partial-control results on the WLASL test split. Higher is better for accuracy, localized diversity, and hypervolume; lower is better for angle. Values average three training seeds; bold marks the best 26-symbol method, and paired 95% intervals are reported in text.

Method	Acc.(%) $\uparrow$	Ang. $^{\circ}$ $\downarrow$	LocDiv $^{\circ}$ $\uparrow$	HV $\uparrow$
<i>Information-rich controls</i>				
Raster skeleton	56.9	25.13	12.82	9.64
Continuous Gauss.	51.3	30.15	12.12	8.94
<i>26-symbol controls</i>				
Learned VQ-26	51.5	29.99	11.36	8.87
Hard HL	52.4	29.48	12.60	9.15
SetHL w/o sph. loss	51.3	30.63	12.42	9.10
SetHL-Wan	52.9	29.04	11.71	8.59

SetHL obtains 8.59 hypervolume, versus 8.94 for continuous and 8.87 for VQ-26. The paired SetHL–continuous difference in the primary endpoint is -0.35 (95% CI [-1.06, 0.34]); against VQ-26 it is -0.29 ([-0.80, 0.18]). The predeclared HV-superiority criterion is therefore not met; SetHL nevertheless has the best accuracy and angular error among the matched 26-symbol controllers. Against the no-spherical-loss ablation, SetHL changes accuracy by 1.56 points ([0.51, 2.65]), angle by -1.59° ([-3.07, -0.43]), and HV by -0.52 ([-0.88, -0.14]). Center-only and hard-HL inference reach 8.86 and **9.15** HV, respectively, confirming that spherical geometry favors semantic compliance rather than uniformly increasing diversity. Relative to continuous control, SetHL changes specified accuracy by 1.59 percentage points ([0.47, 2.79]), angular error by -1.11° ([-2.33, -0.04]), leakage by -3.87 points ([-6.05, -1.72]), and detection by 3.47 points ([1.94, 5.06]); its localized-diversity difference is -0.41° ([-1.80, 1.04]). Against learned VQ-26, SetHL changes specified accuracy by 1.34 percentage points ([0.12, 2.63]) and leakage by -2.80 points ([-4.97, -0.77]). A validation-selected whole-finger mask reproduces the semantic advantage over continuous control: accuracy changes by 2.48 points ([0.97, 4.16]), angle by -2.15° ([-3.75, -0.68]), leakage by -3.74 points ([-5.85, -1.72]), detection by 3.29 points ([1.96, 4.61]), and acceleration error by -0.086 ([-0.165, -0.021]). Its HV change is 0.24 ([-0.01, 0.51]), so frontier superiority remains unsupported. Under full articulation control, SetHL obtains 52.9% HL accuracy versus 51.3% for continuous control (paired difference 1.64 points, 95% CI [0.42, 2.98]). On all 269 test clips, interval-conditioned denoising MSE is 0.1023 for SetHL and 0.1053 for continuous control (paired change -0.0030, 95% CI [-0.0035, -0.0024]). R3D feature similarity is 0.938 versus 0.931 (paired change 0.007, 95% CI [0.005, 0.009]), indicating no loss of source-level video consistency alongside the semantic gains.

Representation analysis. On held-out directions, fixed HL has 16.91° mean quantization error, versus 16.16° for

VQ-26. Under coordinate noise $\sigma = 0.005$, however, HL retains 46.3% of symbols compared with 45.6% for VQ. This separates geometric precision from stability and tests whether a learned vocabulary alone explains the result. Standalone masked completion favors the continuous model ($17.96^{\circ}/61.8\%$) over SetHL ($24.47^{\circ}/54.7\%$); the spherical term nevertheless reduces SetHL completion error by 0.57° . Thus the video comparison does not conflate set-valued control with nearest-trajectory imputation.

5. CONCLUSION

Partial hand instructions should define constraints, not prematurely collapse generation to one completed skeleton. SetHL extends a fixed, interpretable notation into a distribution-valued controller for Wan. Its value is tested by where variation occurs, not merely by whether generated motion is diverse. The present evaluation is limited to short, hand-focused WLASL clips and tracker-derived pseudo-labels; broader domains and human judgments remain future work.

6. COMPLIANCE WITH ETHICAL STANDARDS

This study used only the publicly available WLASL dataset and involved no new data collection or interaction with human participants; no ethical approval was required.

7. ACKNOWLEDGMENT

OpenAI Codex assisted with language drafting throughout the manuscript, implementation of evaluation and statistical-analysis code, and scripting of Figs. 1–2. The authors reviewed all outputs and take responsibility for the study design, paper content, and reported results.

8. REFERENCES

- [1] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala, “Adding conditional control to text-to-image diffusion models,” in *Proc. IEEE/CVF ICCV*, 2023, pp. 3836–3847.
- [2] Yuwei Guo, Ceyuan Yang, Anyi Rao, Maneesh Agrawala, Dahua Lin, and Bo Dai, “SparseCtrl: Adding sparse controls to text-to-video diffusion models,” *arXiv preprint arXiv:2311.16933*, 2023.
- [3] Xiaoyu Shi, Zhaoyang Huang, Fu-Yun Wang, et al., “Motion-I2V: Consistent and controllable image-to-video generation with explicit motion modeling,” in *Proc. ACM SIGGRAPH*, 2024.
- [4] Zhouxia Wang, Ziyang Yuan, Xintao Wang, et al., “Motionctrl: A unified and flexible motion controller for

- video generation,” in *Proc. ACM SIGGRAPH*, 2024, pp. 1–11.
- [5] Ling Li, WenRui Yang, Xinchun Yu, Junliang Xing, and Xiao-Ping Zhang, “Translating motion to notation: Hand labanotation for intuitive and comprehensive hand movement documentation,” in *Proc. ACM Multimedia*, 2024, pp. 4092–4100.
- [6] Wan Team, “Wan: Open and advanced large-scale video generative models,” *arXiv preprint arXiv:2503.20314*, 2025.
- [7] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen, “Lora: Low-rank adaptation of large language models,” in *Proc. ICLR*, 2022.
- [8] Longbin Ji, Lei Zhong, Pengfei Wei, and Changjian Li, “PoseTraj: Pose-aware trajectory control in video diffusion,” in *Proc. IEEE/CVF CVPR*, 2025, pp. 22776–22785.
- [9] Ruihang Chu, Yefei He, Zhekai Chen, et al., “Wan-Move: Motion-controllable video generation via latent trajectory guidance,” in *Proc. NeurIPS*, 2025.
- [10] Daniel Geng, Charles Herrmann, Junhwa Hur, et al., “Motion prompting: Controlling video generation with motion trajectories,” in *Proc. IEEE/CVF CVPR*, 2025.
- [11] Chenyangguang Zhang, Botao Ye, Boqi Chen, et al., “Controllable egocentric video generation via occlusion-aware sparse 3d hand joints,” *arXiv preprint arXiv:2603.11755*, 2026.
- [12] Yanbo Ding, Xirui Hu, Zhizhi Guo, et al., “MTVCraft: Tokenizing 4d motion for arbitrary character animation,” in *Proc. ICLR*, 2026.
- [13] Jingyun Liang, Min Wei, Shikai Li, et al., “Towards 3d-aware video diffusion models: Render-free human motion control with mesh tokenization,” *arXiv preprint arXiv:2606.02000*, 2026.
- [14] Dongxu Li, Cristian Rodriguez Opazo, Xin Yu, and Hongdong Li, “Word-level deep sign language recognition from video: A new large-scale dataset and methods comparison,” in *Proc. IEEE/CVF WACV*, 2020, pp. 1459–1469.
- [15] Anton Pelykh, Edward Fish, Ozge Mercanoglu Sincan, and Richard Bowden, “Signrefine: Adapting foundational video models for sign language generation,” *arXiv preprint arXiv:2609.08496*, 2026.
- [16] Cong Wang, Zexuan Deng, Zhiwei Jiang, et al., “Advanced sign language video generation with compressed and quantized multi-condition tokenization,” *arXiv preprint arXiv:2506.15980*, 2025.
- [17] Fan Zhang, Valentin Bazarevsky, Andrey Vakunov, et al., “MediaPipe Hands: On-device real-time hand tracking,” *arXiv preprint arXiv:2006.10214*, 2020.
- [18] Du Tran, Heng Wang, Lorenzo Torresani, Jamie Ray, Yann LeCun, and Manohar Paluri, “A closer look at spatiotemporal convolutions for action recognition,” in *Proc. IEEE/CVF CVPR*, 2018, pp. 6450–6459.